

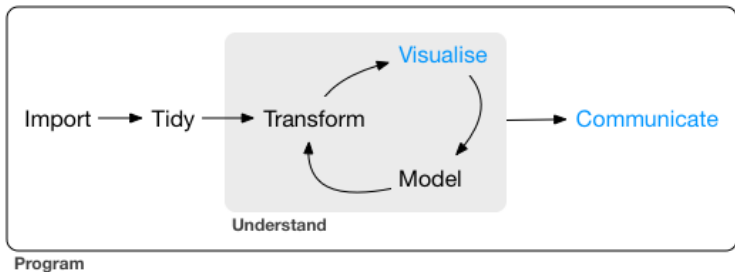
GR5206: lecture 6

*Computational Statistics
And Introduction to Data Science*

Thibault Vatter

Department of Statistics, Columbia University

10/11/2019




Most of the material (e.g., the picture above) is borrowed from

R for data science

- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems
- 5 The layered grammar of graphics
- 6 Labels, axes, annotations and legends
- 7 Colors, zooming and themes

- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems
- 5 The layered grammar of graphics
- 6 Labels, axes, annotations and legends
- 7 Colors, zooming and themes

- Makes no sense to use graphs for very small amounts of data.
- The human brain is capable of grasping a few values.



source: [talkwalker.com](https://www.talkwalker.com)

- Graphs are only as good as the data they display.
- No creativity can produce a good graph from poor data.



- Leinweber (author of *Nerds on Wall Street*):
 - ▶ The S&P500 could be “predicted” at 75% by the butter production in Bangladesh.
 - ▶ ... Or 99% when adding cheese production in the USA, and the population of sheep.

- Graphs shouldn't be more complex than the data they portray.
- Unnecessary complexity can be introduced by irrelevant
 - ▶ decoration
 - ▶ color
 - ▶ 3d effects
- ... Collectively known as “chartjunk”!





Figure 1. Classification of TFBS Regions
TFBS regions for Sp1, cMyc, and p53 were classified based upon proximity to annotations (RefSeq, Sanger hand-curated annotations, GenBank full-length mRNAs, and Ensembl predicted genes). The proximity was calculated from the center of each TFBS region. TFBS regions were classified as follows: within 5 kb of the 5' most exon of a gene, within 5 kb of the 3' terminal exon, or within a gene, novel or outside of any annotation, and pseudogene/ambiguous (TFBS overlapping or flanking pseudogene annotations, limited to chromosome 22, or TFBS regions falling into more than one of the above categories).


source: Cawley S, et al. (2004), Cell 116:499-509, Figure 1

- Graphs shouldn't be distorted pictures of the portrayed values:
 - ▶ Can be either deliberate or accidental.
 - ▶ Useful to know how to produce truth bending graphs.
 - ▶ Misleading often used as a synonym of distorted.



source: statisticshowto.com/misleading-graphs/

- Common sources of distortion:
 - ▶ 3 dimensional “effects”.
 - ▶ linear scaling when using area or volume to represent values.
- The “lie factor”:
 - ▶ Measure of the amount of distortion in a graph.
 - $\text{lie factor} = \frac{\text{size of effect shown in graphic}}{\text{size of effect shown in data}}$
 - Don't take this too seriously.
 - Defined by Ed Tufte of Yale.
 - ▶ If lie factor is > 1 , the graph is exaggerating the effect.



- The three main rules:
 - ▶ If the “story” is simple, keep it simple.
 - ▶ If the “story” is complex, make it look simple.
 - ▶ Tell the truth – do not distort the data.
- Specifically:
 - ▶ There should be a high data to chart ratio.
 - ▶ Use the appropriate graph for the appropriate purpose.
 - Most graphs presented in Excel are POOR CHOICES!
 - In particular, never use a pie chart!
 - ▶ Make sure that the graph is complete:
 - All axes must be labeled.
 - The units should be indicated.
 - There should be a title.
 - A legend can provide needed additional information (e.g., for colors or line types).

“A grammar of graphics is a tool that enables us to concisely describe the components of a graphic. Such a grammar allows us to move beyond named graphics (e.g., the “scatterplot”) and gain insight into the deep structure that underlies statistical graphics.” — Hadley Wickham

- `ggplot2` is an R implementation of the concept:
 - ▶ A coherent system for describing and creating graphs.
 - ▶ Based on [The Grammar of Graphics](#).
 - ▶ Learn one system and apply it in many places.
 - ▶ The equivalent of `dplyr` for graphs.
- To learn more, read [The Layered Grammar of Graphics](#).
- Implementations exist in other languages (e.g., Python)

■ Data from the US EPA on 38 models of car:

```
mpg %>% print(n = 5)
#> # A tibble: 234 x 11
#>   manufacturer model displ  year   cyl trans drv   cty   hwy fl
#>   <chr>          <chr> <dbl> <int> <int> <chr> <chr> <int> <int> <chr>
#> 1 audi          a4      1.8  1999     4 auto~ f    18    29 p
#> 2 audi          a4      1.8  1999     4 manu~ f    21    29 p
#> 3 audi          a4      2    2008     4 manu~ f    20    31 p
#> 4 audi          a4      2    2008     4 auto~ f    21    30 p
#> 5 audi          a4      2.8  1999     6 auto~ f    16    26 p
#> # ... with 229 more rows, and 1 more variable: class <chr>
```

■ Among the variables in mpg are:


- ▶ displ, a car's engine size, in litres.
- ▶ hwy, a car's fuel efficiency on the highway (in miles per gallon).

■ A few questions

- ▶ Do cars with big engines use more fuel ?
- ▶ What does the relationship between engine size and fuel efficiency look like? Positive? Negative? Linear? Nonlinear?

Creating a plot

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy))
```




A graphing template


```
ggplot(data = <DATA>) +  
  <GEOM_FUNCTION>(mapping = aes(<MAPPINGS>))
```

- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems
- 5 The layered grammar of graphics
- 6 Labels, axes, annotations and legends
- 7 Colors, zooming and themes

“The greatest value of a picture is when it forces us to notice what we never expected to see.” — John Tukey




- How to add a third variable to a two dimensional scatterplot?
- By mapping it to an **aesthetic**:
 - ▶ A visual property of the objects in your plot.
 - ▶ Include the size, the shape, or the color of the points.
- We use the words
 - ▶ “**value**” to describe data,
 - ▶ and “**level**” to describe aesthetic properties.



Adding classes to your plot


```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, color = class))
```



- If you prefer British English, use `colour` instead of `color`.


The size aesthetic

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, size = class))  
#> Warning: Using size for a discrete variable is not advised.
```




The alpha aesthetic

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, alpha = class))
```



The shape aesthetic

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, shape = class))
```




class

- 2seater
- ▲ compact
- midsize
- + minivan
- pickup
- * subcompact
- suv

Set the aesthetics manually

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy), color = "blue")
```



- Need values that make sense for that aesthetic:
 - ▶ The name of a color as a character string.
 - ▶ The size of a point in mm.
 - ▶ The shape of a point as a number.





Figure 1: The hollow shapes (0–14) have a border determined by 'color'; the solid shapes (15–18) are filled with 'color'; the filled shapes (21–24) have a border of 'color' and are filled with 'fill'.

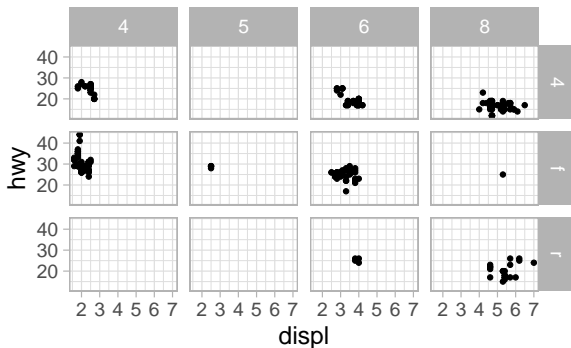
Facets wrap

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy)) +  
  facet_wrap(~ class, nrow = 2)
```



Facets grid



```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy)) +  
  facet_grid(drv ~ cyl)
```



- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations**
- 4 Coordinate systems
- 5 The layered grammar of graphics
- 6 Labels, axes, annotations and legends
- 7 Colors, zooming and themes


How are these two plots similar?

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy))  
  
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy))
```



The linetype aesthetic


```
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy, linetype = drv))
```



- A **geom**:
 - ▶ The object that a plot uses to represent data.
 - ▶ Plots often described by the geom type:
 - Bar charts use bar geoms.
 - Line charts use line geoms.
 - Boxplots use boxplot geoms.
 - ▶ An exception:
 - Scatterplots use the point geom.
- Every **geom** function takes a mapping argument.
- But **not every aesthetic works with every geom**:
 - ▶ **shape** exists for `geom_point` but not for `geom_line`,
 - ▶ and conversely for **linetype**.

Combining two geoms


```
ggplot(data = mpg, mapping = aes(x = displ, y = hwy, color = drv)) +  
  geom_point() +  
  geom_smooth(mapping = aes(linetype = drv))
```




- ggplot2 provides over 30 geoms.
- [extension packages](#) provide even more.
- Use [RStudio's data visualization cheatsheet](#).
- To learn more about any single geom, use help:
 `?geom_smooth`.

Geoms and legends

```
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy))
```





```
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy, group = drv))
```




Geoms and legends

```
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy, color = drv))
```




Multiple geoms in the same plot

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy)) +  
  geom_smooth(mapping = aes(x = displ, y = hwy))
```




A better way

```
ggplot(data = mpg, mapping = aes(x = displ, y = hwy)) +  
  geom_point() +  
  geom_smooth()
```




Local vs global mappings

```
ggplot(data = mpg, mapping = aes(x = displ, y = hwy)) +  
  geom_point(mapping = aes(color = class)) +  
  geom_smooth()
```




Layer dependent data

```
ggplot(data = mpg, mapping = aes(x = displ, y = hwy)) +  
  geom_point(mapping = aes(color = class)) +  
  geom_smooth(data = filter(mpg, class == "subcompact"), se = FALSE)
```



- The diamonds dataset:
 - ▶ About 54,000 diamonds.
 - ▶ Information about price, carat, color, clarity, and cut for each.

```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut))
```



- Other graphs, like bar charts, calculate new values to plot.
 - ▶ Bar charts, histograms, and frequency polygons:
 - Bin data.
 - Plot bin counts (number of points falling in each bin).
 - ▶ Smoothers:
 - Fit a model to your data.
 - Plot predictions from the model.
 - ▶ Boxplots:
 - Compute a robust summary of the distribution.
 - Display a specially formatted box.

■ A stat:

- ▶ The algorithm used to calculate new values for a graph.
- ▶ Short for statistical transformation.


1. `geom_bar()` begins with the **diamonds** data set

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
...

2. `geom_bar()` transforms the data with the "count" stat, which returns a data set of cut values and counts.

cut	count	prop
Fair	1610	1
Good	4906	1
Very Good	12082	1
Premium	13791	1
Ideal	21551	1

3. `geom_bar()` uses the transformed data to build the plot. cut is mapped to the x axis, count is mapped to the y axis.



- `ggplot2` provides over 20 stats.
- Each stat is a function, get help as usual, e.g. `?stat_bin`.
- Use [RStudio's data visualization cheatsheet](#) for a complete list.

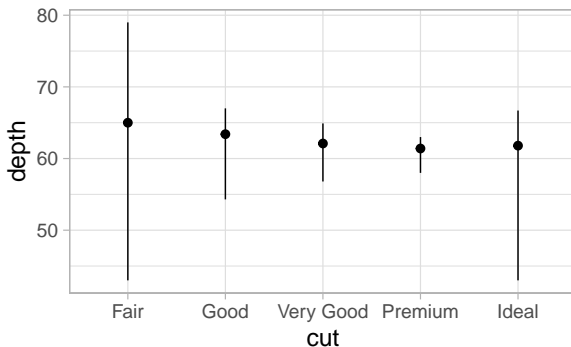
- Every geom has a default stat and conversely.
 - ▶ ?geom_bar shows that the default value for stat is “count”.
 - ▶ Means that geom_bar() uses stat_count().
 - ▶ ?stat_count has a section called “Computed variables” with two new variables: count and prop.
- You can generally use geoms and stats interchangeably!

```
ggplot(data = diamonds) +  
  stat_count(mapping = aes(x = cut))
```

- Typically, use geoms without worrying about the stat.
- Three reasons to use a stat explicitly:
 - ▶ To override the default stat.
 - ▶ To override the default mapping from transformed variables to aesthetics.
 - ▶ To draw greater attention to the stat in your code.


Use a stat explicitly I

```
ggplot(data = diamonds) +  
  stat_summary(  
    mapping = aes(x = cut, y = depth),  
    fun.ymin = min,  
    fun.ymax = max,  
    fun.y = median  
  )
```




Use a stat explicitly II

```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, y = ..prop.., group = 1))
```




Use a stat explicitly III

```
demo <- tribble(~cut,      ~freq,  
               "Fair",    1610,  
               "Good",    4906,  
               "Very Good", 12082,  
               "Premium", 13791,  
               "Ideal",   21551)  
ggplot(data = demo) +  
  geom_bar(mapping = aes(x = cut, y = freq), stat = "identity")
```




The fill aesthetic

```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = cut))
```



Fill and position adjustments


```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = clarity))
```



- Automatically stacked by the **position adjustment**.
- ?position_stack to learn more.

Fill with position = "identity"


```
ggplot(data = diamonds, mapping = aes(x = cut, fill = clarity)) +  
  geom_bar(alpha = 1/5, position = "identity")
```



- Not very useful for bars because of overlap.
- ?position_identity to learn more.

Fill with position = "fill"


```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = clarity), position = "fill")
```



- Makes it easier to compare proportions across groups.
- ?position_fill to learn more.

Fill with position = "dodge"


```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = clarity), position = "dodge")
```



- Makes it easier to compare individual values.
- ?position_dodge to learn more.

position = "jitter"

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_point() +  
  geom_point(position = "jitter", color = "red")
```





- Graph less/**more** accurate/**revealing** at small/**large** scales.
- ?position_jitter to learn more.

- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems**
- 5 The layered grammar of graphics
- 6 Labels, axes, annotations and legends
- 7 Colors, zooming and themes

- The most complicated part of `ggplot2`.
- Default: the Cartesian coordinate system.
- Other systems occasionally helpful:
 - ▶ `coord_flip()` switches the x and y axes.
 - ▶ `coord_quickmap()` sets the aspect ratio correctly for maps.
 - ▶ `coord_polar()` uses polar coordinates.



```
ggplot(data = mpg, mapping = aes(x = class, y = hwy)) +  
  geom_boxplot()  
ggplot(data = mpg, mapping = aes(x = class, y = hwy)) +  
  geom_boxplot() +  
  coord_flip()
```



- Useful for:
 - ▶ horizontal boxplots,
 - ▶ and long labels.

coord_polar()

```
bar <- ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = cut),  
            show.legend = FALSE, width = 1) +  
  theme(aspect.ratio = 1) + labs(x = NULL, y = NULL)  
bar + coord_flip()  
bar + coord_polar()
```



- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems
- 5 The layered grammar of graphics**
- 6 Labels, axes, annotations and legends
- 7 Colors, zooming and themes


```
ggplot(data = <DATA>) +  
  <GEOM_FUNCTION>(  
    mapping = aes(<MAPPINGS>),  
    stat = <STAT>,  
    position = <POSITION>  
  ) +  
  <COORDINATE_FUNCTION> +  
  <FACET_FUNCTION>
```

- A formal system for building plots,
- Uniquely describes *any* plot as a combination of
 - ▶ a dataset,
 - ▶ a geom,
 - ▶ a set of mappings,
 - ▶ a stat,
 - ▶ a position adjustment,
 - ▶ a coordinate system,
 - ▶ and a faceting scheme.

Example

1. Begin with the **diamonds** data set

2. Compute counts for each cut value with **stat_count()**.

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
...

 **stat_count()**

cut	count	prop
Fair	1610	1
Good	4906	1
Very Good	12082	1
Premium	13791	1
Ideal	21551	1

Example

3. Represent each observation with a bar.
4. Map the **fill** of each bar to the **..count..** variable.

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
...

stat_count()

cut	count	prop
Fair	1610	1
Good	4906	1
Very Good	12082	1
Premium	13791	1
Ideal	21551	1

Example


5. Place geoms in a cartesian coordinate system.

6. Map the y values to `..count..` and the x values to `cut`.

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
...


stat_count()

cut	count	prop
Fair	1610	1
Good	4906	1
Very Good	12082	1
Premium	13791	1
Ideal	21551	1



- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems
- 5 The layered grammar of graphics
- 6 Labels, axes, annotations and legends**
- 7 Colors, zooming and themes


```
ggplot(mpg, aes(displ, hwy)) + geom_point(aes(color = class)) +  
  geom_smooth(se = FALSE) +  
  labs(title = "Fuel efficiency decreases with engine size")
```




- Avoid titles that just describe what the plot is!

- subtitle: additional details beneath the title.
- caption: text at the bottom right of the plot.


```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = class)) + geom_smooth(se = FALSE) +  
  labs(title = "Fuel efficiency decreases with engine size",  
        subtitle = "Two seaters (sports cars) are an exception  
        because of their light weight",  
        caption = "Data from fueleconomy.gov")
```



```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = class)) +  
  geom_smooth(se = FALSE) +  
  labs(x = "Engine displacement (L)",  
       y = "Highway fuel economy (mpg)",  
       color = "Car type")
```





```
df <- tibble(x = runif(10),  
             y = runif(10))  
ggplot(df, aes(x, y)) + geom_point() +  
  labs(x = quote(sum(x[i] ^ 2, i == 1, n)),  
       y = quote(alpha + beta + frac(delta, theta)))
```




Annotations

```
best_in_class <- mpg %>%  
  group_by(class) %>%  
  filter(row_number(desc(hwy)) == 1)  
  
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = class)) +  
  geom_text(aes(label = model), data = best_in_class)
```




■ Use the **ggrepel** package!

```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = class)) +  
  geom_point(size = 3, shape = 1, data = best_in_class) +  
  ggrepel::geom_label_repel(aes(label = model), data = best_in_class)
```




Replace legend by labels on the plot

```
class_avg <- mpg %>% group_by(class) %>%  
  summarise(displ = median(displ), hwy = median(hwy))  
  
ggplot(mpg, aes(displ, hwy, color = class)) +  
  ggrepel::geom_label_repel(aes(label = class), data = class_avg,  
    size = 6, label.size = 0,  
    segment.color = NA) +  
  geom_point() + theme(legend.position = "none")
```



To add a single label to the plot


```
label <- mpg %>%  
  summarise(displ = max(displ), hwy = max(hwy),  
            label = "Increasing engine size is  
            related to decreasing fuel economy.")  
  
ggplot(mpg, aes(displ, hwy)) + geom_point() +  
  geom_text(aes(label = label), data = label,  
            vjust = "top", hjust = "right")
```



An alternative

```
label <- tibble(displ = Inf, hwy = Inf,  
               label = "Increasing engine size is  
               related to decreasing fuel economy.")
```


```
ggplot(mpg, aes(displ, hwy)) + geom_point() +  
  geom_text(aes(label = label), data = label,  
           vjust = "top", hjust = "right")
```



To automatically add line breaks

```
"Increasing engine size is related to decreasing fuel economy." %>%  
  stringr::str_wrap(width = 40) %>%  
  writeLines()  
#> Increasing engine size is related to  
#> decreasing fuel economy.
```


To control the alignment of the label



- `geom_hline()` and `geom_vline()`:
 - ▶ Add reference lines.
 - ▶ Using e.g. `size = 2` is often a good idea.
- `geom_rect()`:
 - ▶ Draw a rectangle around points of interest.
 - ▶ Boundaries defined by `xmin`, `xmax`, `ymin`, `ymax`.
- `geom_segment()` with the `arrow` argument:
 - ▶ Draw attention to a point with an arrow.
 - ▶ `x/xend` and `y/yend` define the start/end locations.
- The only limit is your imagination (and patience)!


```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = class))
```

```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = class)) +  
  scale_x_continuous() +  
  scale_y_continuous() +  
  scale_color_discrete()
```



class

- 2seater
- compact
- midsize
- minivan
- pickup
- subcompact
- suv




class

- 2seater
- compact
- midsize
- minivan
- pickup
- subcompact
- suv

- To control the ticks on the axes and the keys on the legend:
 - ▶ `breaks`: controls the position of the ticks, or the values associated with the keys.
 - ▶ `labels`: controls the text label associated with each tick/key.


```
ggplot(mpg, aes(displ, hwy)) + geom_point() +  
  scale_y_continuous(breaks = seq(15, 40, by = 5))
```



Axis ticks and legend keys II

- A useful trick for maps, or for publishing plots where you can't share the absolute numbers:


```
ggplot(mpg, aes(displ, hwy)) + geom_point() +  
  scale_x_continuous(labels = NULL) +  
  scale_y_continuous(labels = NULL)
```



- Collectively axes and legends are called **guides**:
 - ▶ Axes are used for x and y aesthetics.
 - ▶ Legends are used for everything else.
 - ▶ You can also use `breaks` and `labels` to control the appearance of legends.
- Breaks and labels for date and datetime scales work differently:
 - ▶ `date_labels`: takes a format specification, see `?readr::parse_datetime()`.
 - ▶ `date_breaks`: takes a string like "2 days" or "1 month".


- Allow e.g. to highlight exactly where the observations occur:

```
presidential %>%  
  mutate(id = 33 + row_number()) %>%  
  ggplot(aes(start, id)) + geom_point() +  
    geom_segment(aes(xend = end, yend = id)) +  
    scale_x_date(NULL, breaks = presidential$start, date_labels = "'%y")
```



```
base <- ggplot(mpg, aes(displ, hwy)) + geom_point(aes(color = class))

base + theme(legend.position = "left")
base + theme(legend.position = "top")
base + theme(legend.position = "bottom")
base + theme(legend.position = "right") # the default
```




- `legend.position = "none"` suppresses the display of the legend!

To control individual legends

- Use `guides()`, `guide_legend()` or `guide_colorbar()`:


```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = class)) +  
  geom_smooth(se = FALSE) +  
  theme(legend.position = "bottom") +  
  guides(color = guide_legend(nrow = 1, override.aes = list(size = 4)))
```



class ● 2seater ● compact ● midsize ● minivan ● pickup ● subcompact


How could we improve the scale?

```
ggplot(diamonds, aes(carat, price)) +  
  geom_bin2d()
```




Log-transform the variables

```
ggplot(diamonds, aes(log10(carat), log10(price))) +  
  geom_bin2d()
```



... or simply replace the scale

```
ggplot(diamonds, aes(carat, price)) +  
  geom_bin2d() +  
  scale_x_log10() +  
  scale_y_log10()
```





- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems
- 5 The layered grammar of graphics
- 6 Labels, axes, annotations and legends
- 7 Colors, zooming and themes**

Replacing color scales

```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = drv), size = 3)
```


```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = drv), size = 3) +  
  scale_color_brewer(palette = "Blues")
```



- Color scales come in two variety:
 - ▶ `scale_color_x()` for the color aesthetics (available in UK/US spellings).
 - ▶ `scale_fill_x()` for the fill aesthetics.

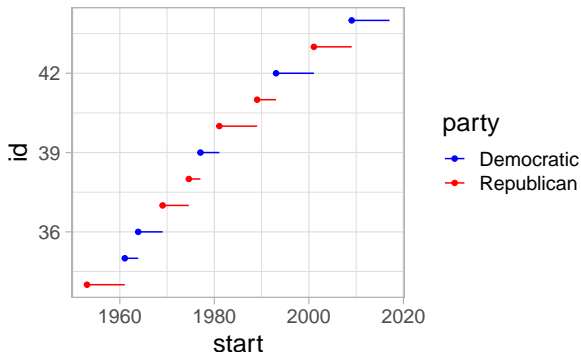
The ColorBrewer scales

- Documented online at <http://colorbrewer2.org/>
- Available via the **RColorBrewer** package.



Using manually defined mappings

```
presidential %>%  
  mutate(id = 33 + row_number()) %>%  
  ggplot(aes(start, id, color = party)) +  
    geom_point() +  
    geom_segment(aes(xend = end, yend = id)) +  
    scale_color_manual(values = c(Republican = "red",  
                                   Democratic = "blue"))
```





Continuous vs diverging color scales

```
df <- data.frame(x = runif(100), y = runif(100),  
                 z1 = rnorm(100), z2 = abs(rnorm(100)))
```

```
ggplot(df, aes(x, y)) +  
  geom_point(aes(color = z2), size = 3)
```

```
ggplot(df, aes(x, y)) +  
  geom_point(aes(color = z1), size = 3) +  
  scale_color_gradient2()
```





A continuous analog of ColorBrewer

■ The viridis package!

```
df <- tibble(x = rnorm(10000), y = rnorm(10000))
```

```
ggplot(df, aes(x, y)) +  
  geom_hex() +  
  coord_fixed()
```



```
ggplot(df, aes(x, y)) +  
  geom_hex() +  
  coord_fixed() +  
  viridis::scale_fill_viridis()
```



■ Three methods:



- ▶ Adjust what data are plotted.
- ▶ Set `xlim` and `ylim` in `coord_cartesian()`.
- ▶ Set the limits in each scale.

```
mpg %>%  
  filter(displ >= 5, displ <= 7, hwy >= 10, hwy <= 30) %>%  
  ggplot(aes(displ, hwy)) +  
    geom_point(aes(color = class)) + geom_smooth()  
  
ggplot(mpg, mapping = aes(displ, hwy)) +  
  geom_point(aes(color = class)) + geom_smooth() +  
  coord_cartesian(xlim = c(5, 7), ylim = c(10, 30))
```



Zooming cont'd



```
suv <- mpg %>%  
  filter(class == "suv")  
compact <- mpg %>%  
  filter(class == "compact")  
  
ggplot(suv, aes(displ, hwy, color = drv)) +  
  geom_point()  
ggplot(compact, aes(displ, hwy, color = drv)) +  
  geom_point()
```




- Training the scales with the limits of the full data:


```
x_scale <- scale_x_continuous(limits = range(mpg$displ))
y_scale <- scale_y_continuous(limits = range(mpg$hwy))
col_scale <- scale_color_discrete(limits = unique(mpg$drv))

ggplot(suv, aes(displ, hwy, color = drv)) + geom_point() +
  x_scale + y_scale + col_scale
ggplot(compact, aes(displ, hwy, color = drv)) + geom_point() +
  x_scale + y_scale + col_scale
```



```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(aes(color = class)) +  
  geom_smooth(se = FALSE) +  
  theme_dark()
```





- More in add-on packages like [ggthemes](#)!