

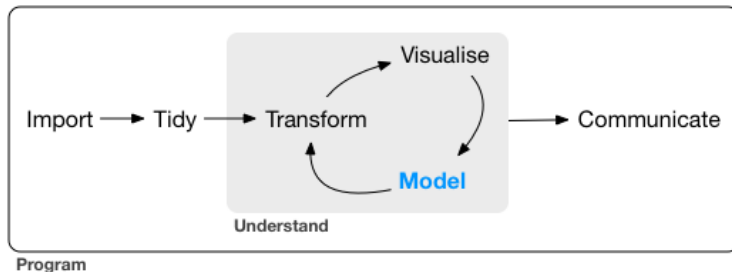
GR5206: lecture 8

*Computational Statistics
And Introduction to Data Science*

Thibault Vatter

Department of Statistics, Columbia University

11/01/2019



- First:
 - ▶ how models work mechanistically (focus on linear models),
 - ▶ how to use models to find patterns in real data.
- Then:
 - ▶ how to use **many** simple models,
 - ▶ how to combine modeling and programming tools.
- As usual, material borrowed from [R for data science](#).

- 1 Model basics
- 2 Model building
- 3 Many models
- 4 List-columns

1 Model basics

2 Model building

3 Many models

4 List-columns

- Each observation can either be used for exploration **OR** confirmation, not both.
- You can use an observation
 - ▶ as many times as you like for exploration,
 - ▶ only once for confirmation.
- When using an observation twice, switch from confirmation to exploration.

- Goals:
 - ▶ Provide a simple low-dimensional summary of a dataset.
 - ▶ Often partition data into patterns and residuals.
 - ▶ Help peel back layers of structure (since strong patterns hide subtler trends).
- The two components of a model:
 - ▶ **Family of models:** a precise, but generic, pattern to capture.
 - A straight line, or a quadratic curve.
 - Equations like $y = a_1 * x + a_2$ or $y = a_1 * x^2$ (with x and y known variables and a_1 and a_2 parameters).
 - ▶ **Fitted model:** member of the family that is closest to the data.
 - $y = 3 * x + 7$ or $y = 9 * x^2$.


All models are wrong, but some are useful. —George Box

- A fitted model is “just” the closest model to the data from a family of models.
- The “best” model (according to some criteria):
 - ▶ isn't necessarily a good model,
 - ▶ isn't necessarily “true”.
- **The goal is not to uncover truth, but to discover useful approximations.**

```
library(tidyverse)  
library(modelr)
```

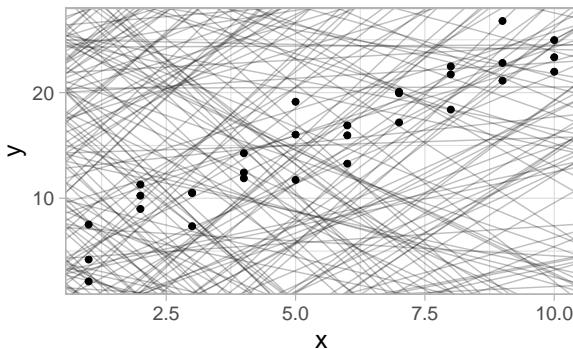
A simulated dataset


```
ggplot(sim1, aes(x, y)) +  
  geom_point(size = 2)
```



Linear family?

```
models <- tibble(a1 = runif(250, -20, 40),  
                 a2 = runif(250, -5, 5))  
  
ggplot(sim1, aes(x, y)) +  
  geom_point(size = 2) +  
  geom_abline(aes(intercept = a1, slope = a2),  
             data = models, alpha = 1/4)
```





- This distance is the difference between
 - ▶ the y value given by the model (the **prediction**),
 - ▶ and the actual y value in the data (the **response**).

■ The model family:

```
model1 <- function(a, data) a[1] + data$x * a[2]
```

```
model1(c(7, 1.5), sim1)
```

```
#> [1] 8.5 8.5 8.5 10.0 10.0 10.0 11.5 11.5 11.5 13.0 13.0 13.0 14.5  
#> [14] 14.5 14.5 16.0 16.0 16.0 17.5 17.5 17.5 19.0 19.0 19.0 20.5 20.5  
#> [27] 20.5 22.0 22.0 22.0
```

■ Root-mean-square error (RMSE):

```
measure_distance <- function(mod, data) {  
  diff <- data$y - model1(mod, data)  
  sqrt(mean(diff ^ 2))  
}
```

```
measure_distance(c(7, 1.5), sim1)
```

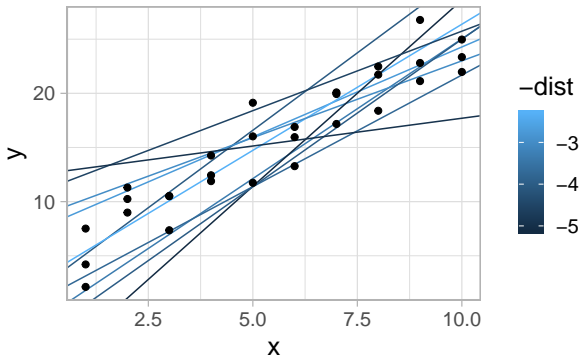
```
#> [1] 2.67
```

```
sim1_dist <- function(a1, a2) measure_distance(c(a1, a2), sim1)

(models <- models %>% mutate(dist = map2_dbl(a1, a2, sim1_dist)))
#> # A tibble: 250 x 3
#>       a1      a2  dist
#>   <dbl>  <dbl> <dbl>
#> 1 -15.2  0.0889  30.8
#> 2  30.1 -0.827   13.2
#> 3  16.0  2.27    13.2
#> 4 -10.6  1.38    18.7
#> 5 -19.6 -1.04    41.8
#> 6   7.98  4.59    19.3
#> 7   9.87 -2.01    20.5
#> 8  -2.61 -4.50    46.9
#> 9  24.0  0.762    13.4
#> 10 26.4 -2.82    14.9
#> # ... with 240 more rows
```

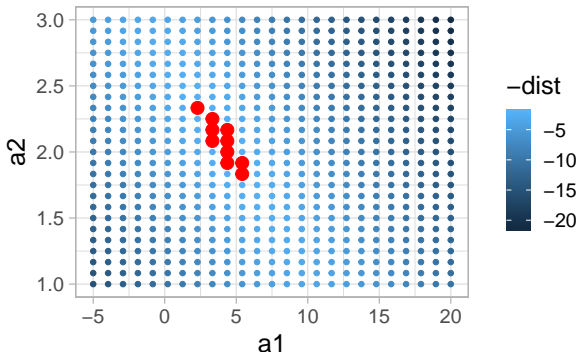
The 10 best models

```
ggplot(sim1, aes(x, y)) +  
  geom_abline(aes(intercept = a1, slope = a2, color = -dist),  
             data = models %>% filter(rank(dist) <= 10)) +  
  geom_point(size = 2)
```



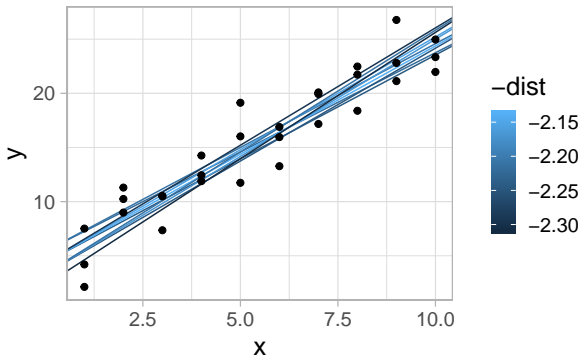
Grid search

```
grid <- expand.grid(a1 = seq(-5, 20, length = 25),  
                  a2 = seq(1, 3, length = 25)) %>%  
  mutate(dist = map2_dbl(a1, a2, sim1_dist))  
  
ggplot(grid, aes(a1, a2)) +  
  geom_point(aes(color = -dist)) +  
  geom_point(data = grid %>% filter(rank(dist) <= 10),  
            size = 4, color = "red")
```



Grid search cont'd

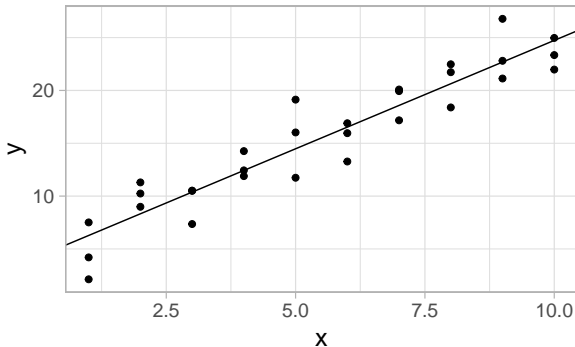
```
ggplot(sim1, aes(x, y)) +  
  geom_abline(aes(intercept = a1, slope = a2, color = -dist),  
             data = grid %>% filter(rank(dist) <= 10)) +  
  geom_point(size = 2)
```



Newton-Raphson search

```
(best <- optim(c(0, 0), measure_distance, data = sim1)$par)
#> [1] 4.22 2.05
```

```
ggplot(sim1, aes(x, y)) +
  geom_point(size = 2) +
  geom_abline(intercept = best[1], slope = best[2])
```



- General form:

- ▶ $y = a_1 + a_2 * x_1 + a_3 * x_2 + \dots + a_n * x_{(n-1)}.$

- `lm()`:

- ▶ Specify the model family using formulas!
 - ▶ E.g., $y \sim x$ is translated to a function like $y = a_1 + a_2 * x.$

```
sim1_mod <- lm(y ~ x, data = sim1)
coef(sim1_mod)
#> (Intercept)          x
#>      4.22      2.05
```

- Two interesting quantities to look at:

- ▶ The **predictions**.
 - ▶ The **residuals**.

- But before that...

- ▶ A statistics digression.
 - ▶ Material borrowed from Prof. Avella-Medina.

- Assume the sample of pairs (\mathbf{X}_i, Y_i) is such that

$$Y_i = \beta_0 + \beta_1 X_{i1} + \beta_1 X_{i2} + \cdots + \beta_d X_{id} + \varepsilon_i, \quad i = 1, \dots, n$$

- Where

- ▶ Response variable: Y_i
- ▶ Covariates: $X_{i1}, X_{i2}, \dots, X_{id}$
- ▶ Noise term: ε_i , assumed to be i.i.d. with $\mathbb{E}[\varepsilon_i] = 0$ and $\text{cov}(X_{ij}, \varepsilon_i) = 0$ for all $j = 1, \dots, d$.

- In matrix form:

$$\underbrace{\begin{bmatrix} Y_1 \\ Y_2 \\ \vdots \\ Y_n \end{bmatrix}}_{\mathbf{Y}} = \underbrace{\begin{bmatrix} 1 & X_{11} & X_{12} & \cdots & X_{1d} \\ 1 & X_{21} & X_{22} & \cdots & X_{2d} \\ \vdots & \vdots & \vdots & & \vdots \\ 1 & X_{n1} & X_{n2} & \cdots & X_{nd} \end{bmatrix}}_{\mathbf{X}} \underbrace{\begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_d \end{bmatrix}}_{\boldsymbol{\beta}} + \underbrace{\begin{bmatrix} \varepsilon_1 \\ \varepsilon_2 \\ \vdots \\ \varepsilon_n \end{bmatrix}}_{\boldsymbol{\varepsilon}}$$

- The least squares estimator (LSE) is defined as

$$\begin{aligned}\hat{\beta} &= \operatorname{argmin}_{\beta \in \mathbb{R}^p} \sum_{i=1}^n (Y_i - \beta_0 - \sum_{j=1}^d X_{ij} \beta_j)^2 = \operatorname{argmin}_{\beta \in \mathbb{R}^p} \sum_{i=1}^n (Y_i - \mathbf{x}_i^T \beta)^2 \\ &= \operatorname{argmin}_{\beta \in \mathbb{R}^p} \{(\mathbf{Y} - \mathbf{X}^T \beta)^T (\mathbf{Y} - \mathbf{X}^T \beta)\} = \operatorname{argmin}_{\beta \in \mathbb{R}^p} \|\mathbf{Y} - \mathbf{X}^T \beta\|_2^2\end{aligned}$$

- If $\operatorname{rank}(\mathbf{X}) = 1 + d = p$, we have a closed form solution

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$


- Hence, under the model assumptions, $\hat{\beta}$ is unbiased since

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T (\mathbf{X} \beta + \varepsilon) = \beta + (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \varepsilon$$

- $\hat{Y} = \mathbf{X}\hat{\beta}$ has the geometric interpretation of being the projection of \mathbf{Y} onto the plane spanned by the columns of:

$$\mathbf{X}\hat{\beta} = \mathbf{X}(\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{Y} = \mathbf{P}\mathbf{Y}$$

- Linear least squares fitting with $\mathbf{X} \in \mathbb{R}^2$. We seek the linear function of \mathbf{X} that minimizes the sum of squared residuals from Y (Friedman, Hastie and Tibshirani 2008).



Assuming we have an i.i.d sample of pairs $(\mathbf{X}_1, Y_1), \dots, (\mathbf{X}_n, Y_n)$ one can establish we can establish theoretical guarantees for $\hat{\beta}$

- Consistency

$$\hat{\beta} \xrightarrow[n \rightarrow \infty]{\mathcal{P}} \beta$$

- Asymptotic normality

$$\sqrt{n}(\hat{\beta} - \beta) \xrightarrow[n \rightarrow \infty]{\mathcal{D}} N(\mathbf{0}, \sigma^2 \mathbf{Q}^{-1}),$$

where $\mathbf{Q} = \mathbb{E}[\mathbf{X}_1 \mathbf{X}_1^T]$.

- Further assuming that the errors ε_i are i.i.d. $N(0, \sigma^2)$
 - ▶ $Y_i | \mathbf{X}_i = \mathbf{x}_i \sim N(\mathbf{x}_i^T \boldsymbol{\beta}, \sigma^2)$.
 - ▶ $\hat{\boldsymbol{\beta}}$ is also the MLE of $\boldsymbol{\beta}$.
 - ▶ For a fixed design matrix \mathbf{X}

$$\hat{\boldsymbol{\beta}} \sim N_p(\boldsymbol{\beta}, \sigma^2(\mathbf{X}^T \mathbf{X})^{-1}).$$

- ▶ $\hat{\sigma}^2 = \frac{1}{n-p} \|\mathbf{Y} - \mathbf{X}\hat{\boldsymbol{\beta}}\|_2^2$ is an unbiased estimator of σ^2 . Indeed, one can show that $(n-p) \frac{\hat{\sigma}^2}{\sigma^2} \sim \chi_{n-p}^2$.
- ▶ $(n-p) \frac{\hat{\sigma}^2}{\sigma^2} \sim \chi_{n-p}^2$
- ▶ $\hat{\boldsymbol{\beta}}$ and $\hat{\sigma}^2$ are independent!
- ▶ Hence

$$(\hat{\beta}_j - \beta_j) / \hat{\sigma}_{\hat{\beta}_j} \sim t_{n-p},$$

where $\hat{\sigma}_{\hat{\beta}_j}^2$ is the j th diagonal element of $\hat{\sigma}^2(\mathbf{X}^T \mathbf{X})^{-1}$.

- Log likelihood function:

$$\ell_n(\beta) = \log L_n(\beta, \sigma^2) = -\frac{1}{2} \left\{ n \log \sigma^2 + \frac{1}{\sigma^2} \sum_{i=1}^n (y_i - \mathbf{x}_i^T \beta)^2 \right\}$$

- Since $\hat{\beta} = \hat{\beta}_{ML}$ we see that

$$\hat{\sigma}_{ML}^2 = \frac{1}{n} \|\mathbf{Y} - \mathbf{X}\hat{\beta}\|_2^2$$

- Hence, log likelihood is maximized at

$$\begin{aligned} \ell_n(\beta) &= \log L_n(\hat{\beta}, \hat{\sigma}_{ML}^2) = -\frac{1}{2} \left\{ n \log \hat{\sigma}_{ML}^2 + n \right\} \\ &= -\frac{1}{2} \left\{ n \log \|\mathbf{Y} - \mathbf{X}\hat{\beta}\|_2^2 + n - n \log n \right\} \end{aligned}$$

- When comparing two nested models it is useful to write

$$\mathbf{Y} = \mathbf{X}\beta + \varepsilon = \begin{bmatrix} \mathbf{X}_1 & \mathbf{X}_2 \end{bmatrix} \begin{bmatrix} \beta_1 \\ \beta_2 \end{bmatrix} + \varepsilon = \mathbf{X}_1\beta_1 + \mathbf{X}_2\beta_2 + \varepsilon,$$

where \mathbf{X}_1 is $n \times q$ and \mathbf{X}_2 is $n \times (p - q)$.

- Likelihood ratio statistic:

$$2\{\ell_n(\hat{\beta}) - \ell_n(\hat{\beta}_1^R)\} = n \log \left(\frac{\|\mathbf{Y} - \mathbf{X}_1^T \hat{\beta}_1^R\|_2^2}{\|\mathbf{Y} - \mathbf{X}^T \hat{\beta}\|_2^2} \right) = n \log \left(1 + \frac{p - q}{n - p} F \right)$$

- F-statistic:

$$F = \frac{n - p}{p - q} \frac{\|\mathbf{Y} - \mathbf{X}^T \hat{\beta}_1^R\|_2^2 - \|\mathbf{Y} - \mathbf{X}^T \hat{\beta}\|_2^2}{\|\mathbf{Y} - \mathbf{X}^T \hat{\beta}\|_2^2} \sim F_{p-q, n-p}$$


```
summary(sim1_mod)
#>
#> Call:
#> lm(formula = y ~ x, data = sim1)
#>
#> Residuals:
#>      Min       1Q   Median       3Q      Max
#> -4.147 -1.520  0.133  1.467  4.652
#>
#> Coefficients:
#>              Estimate Std. Error t value Pr(>|t|)
#> (Intercept)    4.221      0.869    4.86  4.1e-05 ***
#> x              2.052      0.140   14.65  1.2e-14 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
#>
#> Residual standard error: 2.2 on 28 degrees of freedom
#> Multiple R-squared:  0.885, Adjusted R-squared:  0.88
#> F-statistic: 215 on 1 and 28 DF, p-value: 1.17e-14
```

```
sim0_mod <- lm(y ~ 1, data = sim1)
anova(sim0_mod, sim1_mod)
#> Analysis of Variance Table
#>
#> Model 1: y ~ 1
#> Model 2: y ~ x
#>   Res.Df  RSS Df Sum of Sq   F  Pr(>F)
#> 1      29 1178
#> 2      28  136  1    1042 215 1.2e-14 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- ... end of the digression!
- Two interesting quantities to look at:
 - ▶ The **predictions**.
 - ▶ The **residuals**.

■ Step 1


```
(grid <- sim1 %>%  
  modelr::data_grid(x))  
#> # A tibble: 10 x 1  
#>       x  
#>   <int>  
#> 1     1  
#> 2     2  
#> 3     3  
#> 4     4  
#> 5     5  
#> 6     6  
#> 7     7  
#> 8     8  
#> 9     9  
#> 10    10
```

■ Step 2

```
(grid <- grid %>%  
  add_predictions(sim1_mod))  
#> # A tibble: 10 x 2  
#>       x pred  
#>   <int> <dbl>  
#> 1     1  6.27  
#> 2     2  8.32  
#> 3     3 10.4  
#> 4     4 12.4  
#> 5     5 14.5  
#> 6     6 16.5  
#> 7     7 18.6  
#> 8     8 20.6  
#> 9     9 22.7  
#> 10    10 24.7
```

■ Step 3

```
ggplot(sim1, aes(x, y)) +  
  geom_point(size = 2) +  
  geom_line(aes(y = pred), data = grid, size = 1)
```




■ Step 1

```
(sim1 <- sim1 %>%  
  add_residuals(sim1_mod))  
#> # A tibble: 30 x 3  
#>       x     y   resid  
#>   <int> <dbl>   <dbl>  
#> 1     1     4.20 -2.07  
#> 2     1     7.51  1.24  
#> 3     1     2.13 -4.15  
#> 4     2     8.99  0.665  
#> 5     2    10.2   1.92  
#> 6     2    11.3   2.97  
#> 7     3     7.36 -3.02  
#> 8     3    10.5   0.130  
#> 9     3    10.5   0.136  
#> 10    4    12.4   0.00763  
#> # ... with 20 more rows
```


■ Step 2

```
ggplot(sim1, aes(resid)) +  
  geom_freqpoly(binwidth = 0.5)
```



■ Step 3

```
ggplot(sim1, aes(x, resid)) +  
  geom_ref_line(h = 0) +  
  geom_point()
```



- What's that?
 - ▶ A way of getting “special behavior”.
 - ▶ “Capture variables” so they can be interpreted by the function.
 - ▶ Sometimes called “Wilkinson-Rogers notation” from [Symbolic Description of Factorial Models for Analysis of Variance](#)
- Behind the scenes:

```
df <- tribble(~y, ~x1, ~x2,  
             4, 2, 5,  
             5, 1, 6)  
model_matrix(df, y ~ x1)  
#> # A tibble: 2 x 2  
#>   `(Intercept)`      x1  
#>   <dbl> <dbl>  
#> 1      1      2  
#> 2      1      1
```

■ Without intercept:

```
model_matrix(df, y ~ x1 - 1)
#> # A tibble: 2 x 1
#>       x1
#>   <dbl>
#> 1     2
#> 2     1
```

■ Adding a second variable:

```
model_matrix(df, y ~ x1 + x2)
#> # A tibble: 2 x 3
#>   `(Intercept)`    x1    x2
#>   <dbl> <dbl> <dbl>
#> 1         1     2     5
#> 2         1     1     6
```




```
df <- tribble(~ sex, ~ response,  
              "male", 1,  
              "female", 2,  
              "male", 1)  
model_matrix(df, response ~ sex)  
#> # A tibble: 3 x 2  
#>   `(Intercept)` sexmale  
#>       <dbl>    <dbl>  
#> 1           1        1  
#> 2           1        0  
#> 3           1        1
```

- Why doesn't R also create a sexfemale column?

Another simulated dataset

```
ggplot(sim2, aes(x, y)) +  
  geom_point(size = 2)
```



```
mod2 <- lm(y ~ x, data = sim2)
```

```
(grid <- sim2 %>%  
  data_grid(x) %>%  
  add_predictions(mod2))
```

```
#> # A tibble: 4 x 2
```

```
#>   x      pred
```

```
#>   <chr> <dbl>
```

```
#> 1 a      1.15
```


```
#> 2 b      8.12
```

```
#> 3 c      6.13
```

```
#> 4 d      1.91
```

Visualize the results

```
ggplot(sim2, aes(x)) +  
  geom_point(aes(y = y), size = 2) +  
  geom_point(data = grid, aes(y = pred), color = "red", size = 4)
```




What's happening here?

```
tibble(x = "e") %>%  
  add_predictions(mod2)  
#> Error in model.frame.default(Terms, newdata, na.action = na.action,  
#> xlev = object$xlevels): factor x has new level e
```

Interactions (cont. and cat.)

```
ggplot(sim3, aes(x1, y)) +  
  geom_point(aes(color = x2))
```



Two possible models

```
mod1 <- lm(y ~ x1 + x2, data = sim3)
mod2 <- lm(y ~ x1 * x2, data = sim3)
```

■ Note that:

- ▶ $y \sim x1 + x2$ becomes $y = a0 + a1 * x1 + a2 * x2$.
- ▶ $y \sim x1 * x2$ becomes $y = a0 + a1 * x1 + a2 * x2 + a12 * x1 * x2$.

Two new tricks to visualize them

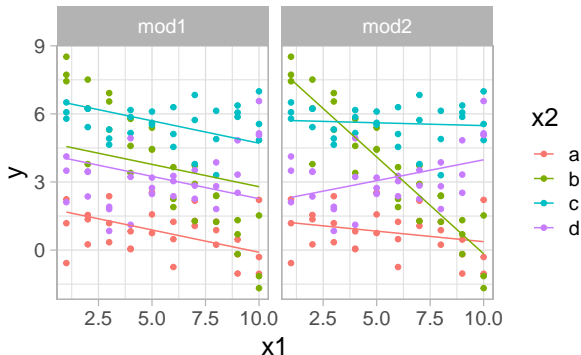
- Give `data_grid()` both variables.
- To generate predictions from both models simultaneously, use
 - ▶ `gather_predictions()` to add predictions as rows,
 - ▶ or `spread_predictions()` to add predictions as columns.

```
(grid <- sim3 %>%  
  data_grid(x1, x2) %>%  
  gather_predictions(mod1, mod2))
```

```
#> # A tibble: 80 x 4  
#>   model    x1 x2    pred  
#>   <chr> <int> <fct> <dbl>  
#> 1 mod1     1 a     1.67  
#> 2 mod1     1 b     4.56  
#> 3 mod1     1 c     6.48  
#> 4 mod1     1 d     4.03  
#> 5 mod1     2 a     1.48  
#> 6 mod1     2 b     4.37  
#> 7 mod1     2 c     6.28  
#> 8 mod1     2 d     3.84  
#> 9 mod1     3 a     1.28  
#> 10 mod1    3 b     4.17  
#> # ... with 70 more rows
```


Using facetting

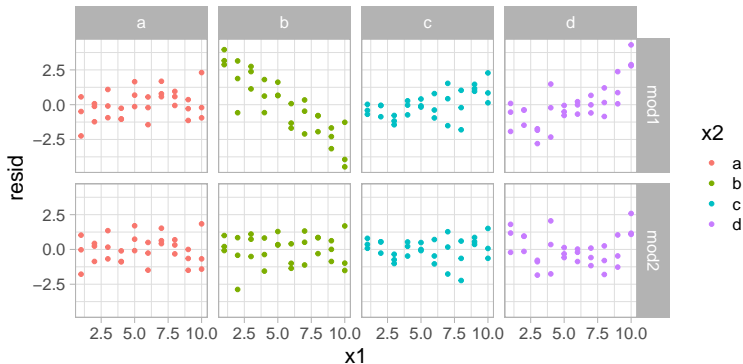
```
ggplot(sim3, aes(x1, y, color = x2)) +  
  geom_point() +  
  geom_line(data = grid, aes(y = pred)) +  
  facet_wrap(~ model)
```



Which model is better?

```
sim3 <- sim3 %>% gather_residuals(mod1, mod2)
```

```
ggplot(sim3, aes(x1, resid, color = x2)) +  
  geom_point() +  
  facet_grid(model ~ x2)
```



Which model is better?

■ Remember slide 23

```
anova(mod1, mod2)
#> Analysis of Variance Table
#>
#> Model 1: y ~ x1 + x2
#> Model 2: y ~ x1 * x2
#>   Res.Df RSS Df Sum of Sq    F Pr(>F)
#> 1      115 270
#> 2      112 118  3      153 48.5 <2e-16 ***
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Interactions (two continuous)


```
mod1 <- lm(y ~ x1 + x2, data = sim4)
mod2 <- lm(y ~ x1 * x2, data = sim4)

(grid <- sim4 %>%
  data_grid(x1 = seq_range(x1, 5), x2 = seq_range(x2, 5)) %>%
  gather_predictions(mod1, mod2))

#> # A tibble: 50 x 4
#>   model    x1    x2   pred
#>   <chr> <dbl> <dbl> <dbl>
#> 1 mod1    -1    -1  0.996
#> 2 mod1    -1   -0.5 -0.395
#> 3 mod1    -1     0 -1.79
#> 4 mod1    -1    0.5 -3.18
#> 5 mod1    -1     1 -4.57
#> 6 mod1   -0.5   -1  1.91
#> 7 mod1   -0.5  -0.5  0.516
#> 8 mod1   -0.5     0 -0.875
#> 9 mod1   -0.5    0.5 -2.27
#> 10 mod1  -0.5     1 -3.66
#> # ... with 40 more rows
```


?seq_range for other arguments (e.g., pretty = TRUE for tables).

```
ggplot(grid, aes(x1, x2)) + geom_tile(aes(fill = pred)) +  
  facet_wrap(~ model)
```




Slices with respect to x2

```
ggplot(grid, aes(x1, pred, color = x2, group = x2)) +  
  geom_line() +  
  facet_wrap(~ model)
```



Slices with respect to x1

```
ggplot(grid, aes(x2, pred, color = x1, group = x1)) +  
  geom_line() +  
  facet_wrap(~ model)
```



Which model is better?

■ Remember slide 23

```
anova(mod1, mod2)
#> Analysis of Variance Table
#>
#> Model 1: y ~ x1 + x2
#> Model 2: y ~ x1 * x2
#>   Res.Df  RSS Df Sum of Sq    F Pr(>F)
#> 1      297 1323
#> 2      296 1278  1      45.2 10.5 0.0014 **
#> ---
#> Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


- $\log(y) \sim \sqrt{x_1} + x_2$ is transformed to $\log(y) = a_1 + a_2 * \sqrt{x_1} + a_3 * x_2$.
- If the transformation involves $+$, $*$, $^$, or $-$, wrap it in $I()$:
 - ▶ $y \sim x + I(x^2) \equiv y = a_1 + a_2 * x + a_3 * x^2$.
 - ▶ $y \sim x^2 + x \equiv y \sim x * x + x \equiv y = a_1 + a_2 * x$.

```
df <- tribble(~y, ~x, 1, 1, 2, 2, 3, 3)
model_matrix(df, y ~ x^2 + x)
model_matrix(df, y ~ I(x^2) + x)
#> # A tibble: 3 x 2
#>   `(Intercept)`      x
#>   <dbl> <dbl>
#> 1         1         1
#> 2         1         2
#> 3         1         3
#> # A tibble: 3 x 3
#>   `(Intercept)` `I(x^2)`      x
#>   <dbl> <dbl> <dbl>
#> 1         1         1         1
#> 2         1         4         2
#> 3         1         9         3
```


- To get $y = a_1 + a_2 * x + a_3 * x^2$:

```
model_matrix(df, y ~ poly(x, 2))  
#> # A tibble: 3 x 3  
#>   `(Intercept)` `poly(x, 2)1` `poly(x, 2)2`  
#>         <dbl>         <dbl>         <dbl>  
#> 1             1      -7.07e- 1         0.408  
#> 2             1      -7.85e-17        -0.816  
#> 3             1       7.07e- 1         0.408
```

```
library(splines)
model_matrix(df, y ~ ns(x, 2))
#> # A tibble: 3 x 3
#>   `(Intercept)` `ns(x, 2)1` `ns(x, 2)2`
#>   <dbl>         <dbl>         <dbl>
#> 1         1         0         0
#> 2         1      0.566     -0.211
#> 3         1      0.344      0.771
```

A non-linear function

```
sim5 <- tibble(x = seq(0, 3.5 * pi, length = 50),  
              y = 4 * sin(x) + rnorm(length(x)))  
  
ggplot(sim5, aes(x, y)) + geom_point()
```




■ Three models using splines:

```
mod1 <- lm(y ~ ns(x, 1), data = sim5)  
mod2 <- lm(y ~ ns(x, 2), data = sim5)  
mod5 <- lm(y ~ ns(x, 5), data = sim5)
```

```
grid <- sim5 %>%
  data_grid(x = seq_range(x, n = 50, expand = 0.1)) %>%
  gather_predictions(mod1, mod2, mod5, .pred = "y")

ggplot(sim5, aes(x, y)) + geom_point() +
  geom_line(data = grid, color = "red") +
  facet_wrap(~ model)
```



```
options(na.action = na.warn)
df <- tribble(~x, ~y,
              1, 2.2,
              2, NA,
              3, 3.5,
              4, 8.3,
              NA, 10)

mod <- lm(y ~ x, data = df)
#> Warning: Dropping 2 rows with missing values
mod <- lm(y ~ x, data = df, na.action = na.exclude)
nobs(mod)
#> [1] 3
```

- **Generalized linear models**, e.g. `stats::glm()`:
 - ▶ While LMs assume continuous responses/Gaussian errors, GLMs extend them to other distributions, including non-continuous responses (e.g. binary data or counts).
- **Generalized additive models**, e.g. `mgcv::gam()`:
 - ▶ Extend GLMs to incorporate smooth functions
 - ▶ Formulas like $y \sim s(x)$ become equations like $y = f(x)$.
- **Penalized linear models**, e.g. `glmnet::glmnet()`:
 - ▶ Add penalties to favor simpler models and “generalize” better.
- **Robust linear models**, e.g. `MASS::rlm()`:
 - ▶ Tweaks distance to downweight outliers.
 - ▶ Less sensitive to outliers, but slightly worse without outliers.
- **Trees**, e.g. `rpart::rpart()`:
 - ▶ Piece-wise constant models splitting the data into small pieces.
 - ▶ Powerful when aggregated as **random forests** (e.g. `randomForest::randomForest()`) or **gradient boosting machines** (e.g. `xgboost::xgboost()`).

1 Model basics

2 Model building




3 Many models

4 List-columns

- To partition data into pattern and residuals:
 - ▶ Find patterns with visualization.
 - ▶ Make them concrete and precise with a model.
 - ▶ Repeat 1. and 2. after replacing the old response variable with the residuals from the model.
- How about large and complex datasets?
 - ▶ ML approaches “simply” focus on predictive ability.
 - ▶ Issues:
 - black boxes,
 - (sometimes) hard to use domain knowledge,
 - (often) difficult to assess whether or not the model will continue to work in the long-term
 - ▶ Usually, a combination of both approaches is preferred.

The diamonds dataset


```
ggplot(diamonds, aes(cut, price)) + geom_boxplot()  
ggplot(diamonds, aes(color, price)) + geom_boxplot()  
ggplot(diamonds, aes(clarity, price)) + geom_boxplot()
```



- Why are low quality diamonds more expensive?

Price and carat


```
ggplot(diamonds, aes(carat, price)) +  
  geom_hex(bins = 50)
```



A couple of tweaks

- Focus on diamonds < 2.5 carats (99.7% of the data).
- Log-transform the carat and price.

```
diamonds2 <- diamonds %>%  
  filter(carat <= 2.5) %>%  
  mutate(lprice = log2(price), lcarat = log2(carat))  
  
ggplot(diamonds2, aes(lcarat, lprice)) +  
  geom_hex(bins = 50)
```




- A simple model:

```
mod_diamond <- lm(lprice ~ lcarat, data = diamonds2)
```


Visualize the predictions

```
grid <- diamonds2 %>%  
  data_grid(carat = seq_range(carat, 20)) %>%  
  mutate(lcarat = log2(carat)) %>%  
  add_predictions(mod_diamond, "lprice") %>%  
  mutate(price = 2 ^ lprice)  
  
ggplot(diamonds2, aes(carat, price)) +  
  geom_hex(bins = 50) +  
  geom_line(data = grid, color = "red", size = 1)
```






Visualize the residuals

```
diamonds2 <- diamonds2 %>%  
  add_residuals(mod_diamond, "lresid")  
  
ggplot(diamonds2, aes(lcarat, lresid)) +  
  geom_hex(bins = 50)
```



Replace price by residuals

```
ggplot(diamonds2, aes(cut, lresid)) + geom_boxplot()  
ggplot(diamonds2, aes(color, lresid)) + geom_boxplot()  
ggplot(diamonds2, aes(clarity, lresid)) + geom_boxplot()
```



A more complicated model


```
mod_diamond2 <- lm(lprice ~ lcarat + color + cut + clarity,
                   data = diamonds2)

(grid <- diamonds2 %>%
  data_grid(cut,
            lcarat = -0.515,
            color = "G",
            clarity = "SI1") %>%
  add_predictions(mod_diamond2))

#> # A tibble: 5 x 5
#>   cut      lcarat color clarity  pred
#>   <ord>    <dbl> <chr>  <chr>   <dbl>
#> 1 Fair      -0.515 G      SI1     11.0
#> 2 Good      -0.515 G      SI1     11.1
#> 3 Very Good -0.515 G      SI1     11.2
#> 4 Premium  -0.515 G      SI1     11.2
#> 5 Ideal     -0.515 G      SI1     11.2
```


Visualize the residuals

```
diamonds2 <- diamonds2 %>%  
  add_residuals(mod_diamond2, "lresid2")  
  
ggplot(diamonds2, aes(lcarat, lresid2)) +  
  geom_hex(bins = 50)
```




The number of daily flights

```
library(nycflights13)
library(lubridate)
daily <- flights %>%
  mutate(date = make_date(year, month, day)) %>%
  group_by(date) %>%
  summarize(n = n())
```

What affects this number?


```
ggplot(daily, aes(date, n)) +  
  geom_line()
```



```
daily <- daily %>% mutate(wday = wday(date, label = TRUE))
```


```
ggplot(daily, aes(wday, n)) + geom_boxplot()
```

```
mod <- lm(n ~ wday, data = daily)
```



Visualize the predictions


```
grid <- daily %>%  
  data_grid(wday) %>%  
  add_predictions(mod, "n")  
  
ggplot(daily, aes(wday, n)) +  
  geom_boxplot() +  
  geom_point(data = grid, color = "red", size = 4)
```



Visualize the residuals


```
daily <- daily %>%  
  add_residuals(mod)
```

```
daily %>%  
  ggplot(aes(date, resid)) +  
  geom_ref_line(h = 0) + geom_line()
```



What happens here?

```
ggplot(daily, aes(date, resid, color = wday)) +  
  geom_ref_line(h = 0) + geom_line()
```




What happens here?

```
daily %>%  
  filter(resid < -100)  
#> # A tibble: 11 x 4  
#>   date          n wday  resid  
#>   <date>      <int> <ord> <dbl>  
#> 1 2013-01-01   842 Tue   -109.  
#> 2 2013-01-20   786 Sun   -105.  
#> 3 2013-05-26   729 Sun   -162.  
#> 4 2013-07-04   737 Thu   -229.  
#> 5 2013-07-05   822 Fri   -145.  
#> 6 2013-09-01   718 Sun   -173.  
#> 7 2013-11-28   634 Thu   -332.  
#> 8 2013-11-29   661 Fri   -306.  
#> 9 2013-12-24   761 Tue   -190.  
#> 10 2013-12-25  719 Wed   -244.  
#> 11 2013-12-31  776 Tue   -175.
```



What happens here?

```
daily %>%  
  ggplot(aes(date, resid)) +  
    geom_ref_line(h = 0) +  
    geom_line(color = "grey50") +  
    geom_smooth(se = FALSE, span = 0.20)  
#> `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Seasonal Saturday effect

```
daily %>%  
  filter(wday == "Sat") %>%  
  ggplot(aes(date, n)) +  
  geom_point() + geom_line() +  
  scale_x_date(NULL, date_breaks = "1 month", date_labels = "%b")
```



- State's school terms: summer break in 2013 was Jun 26–Sep 9.

The three school terms


```
term <- function(date) {  
  cut(date, breaks = ymd(20130101, 20130605, 20130825, 20140101),  
       labels = c("spring", "summer", "fall"))  
}  
daily <- daily %>%  
  mutate(term = term(date))
```

daily

```
#> # A tibble: 365 x 5  
#>   date           n wday   resid term  
#>   <date>       <int> <ord> <dbl> <fct>  
#> 1 2013-01-01   842 Tue   -109. spring  
#> 2 2013-01-02   943 Wed    -19.7 spring  
#> 3 2013-01-03   914 Thu    -51.8 spring  
#> 4 2013-01-04   915 Fri    -52.5 spring  
#> 5 2013-01-05   720 Sat    -24.6 spring  
#> 6 2013-01-06   832 Sun    -59.5 spring  
#> 7 2013-01-07   933 Mon    -41.8 spring  
#> 8 2013-01-08   899 Tue    -52.4 spring  
#> 9 2013-01-09   902 Wed    -60.7 spring  
#> 10 2013-01-10  932 Thu    -33.8 spring  
#> # ... with 355 more rows
```


The three school terms cont'd

```
daily %>%  
  filter(wday == "Sat") %>%  
  ggplot(aes(date, n, color = term)) +  
  geom_point(alpha = 1/3) +  
  geom_line() +  
  scale_x_date(NULL, date_breaks = "1 month", date_labels = "%b")
```



School terms and day of week


```
daily %>%  
  ggplot(aes(wday, n, color = term)) +  
  geom_boxplot()
```



An improved model


```
mod1 <- lm(n ~ wday, data = daily)
mod2 <- lm(n ~ wday * term, data = daily)
```

```
daily %>%
  gather_residuals(without_term = mod1, with_term = mod2) %>%
  ggplot(aes(date, resid, color = model)) +
  geom_line(alpha = 0.75)
```




What's going on here?

```
grid <- daily %>%  
  data_grid(wday, term) %>%  
  add_predictions(mod2, "n")  
  
ggplot(daily, aes(wday, n)) +  
  geom_boxplot() +  
  geom_point(data = grid, color = "red") +  
  facet_wrap(~ term)
```



```
mod3 <- MASS::rlm(n ~ wday * term, data = daily)
```

```
daily %>%  
  add_residuals(mod3, "resid") %>%  
  ggplot(aes(date, resid)) +  
  geom_hline(yintercept = 0, size = 2, color = "white") +  
  geom_line()
```



- Either bundled up into a function:

```
compute_vars <- function(data) {  
  data %>%  
    mutate(term = term(date),  
           wday = wday(date, label = TRUE))  
}
```


- Or directly in the model formula:

```
wday2 <- function(x) wday(x, label = TRUE)  
mod3 <- lm(n ~ wday2(date) * term(date), data = daily)
```

Time of year: an alternative approach COLUMBIA UNIVERSITY IN THE CITY OF NEW YORK

```
library(splines)
mod <- MASS::rlm(n ~ wday * ns(date, 5), data = daily)

daily %>%
  data_grid(wday, date = seq_range(date, n = 13)) %>%
  add_predictions(mod) %>%
  ggplot(aes(date, pred, color = wday)) +
  geom_line() +
  geom_point()
```



- 1 Model basics
- 2 Model building
- 3 Many models**
- 4 List-columns


- To work with large numbers of models, use:
 - ▶ Many simple models to better understand complex datasets.
 - ▶ List-columns to store arbitrary data structures in a data frame.
 - ▶ The **broom** package to turn models into tidy data.
- Note that this part
 - ▶ is harder than the others,
 - ▶ requires a deeper internalization of ideas (e.g., about modeling, data structures, and iteration).

- Summarizes the progression of countries over time using variables like life expectancy and GDP.
- Popularized by Hans Rosling, a Swedish doctor and statistician, in a short video filmed in conjunction with the BBC

```
library(gapminder)
gapminder
#> # A tibble: 1,704 x 6
#>   country      continent  year lifeExp      pop gdpPercap
#>   <fct>        <fct>    <int>   <dbl>    <int>    <dbl>
#> 1 Afghanistan Asia      1952    28.8  8425333    779.
#> 2 Afghanistan Asia      1957    30.3  9240934    821.
#> 3 Afghanistan Asia      1962    32.0 10267083    853.
#> 4 Afghanistan Asia      1967    34.0 11537966    836.
#> 5 Afghanistan Asia      1972    36.1 13079460    740.
#> 6 Afghanistan Asia      1977    38.4 14880372    786.
#> 7 Afghanistan Asia      1982    39.9 12881816    978.
#> 8 Afghanistan Asia      1987    40.8 13867957    852.
#> 9 Afghanistan Asia      1992    41.7 16317921    649.
#> 10 Afghanistan Asia      1997    41.8 22227415    635.
#> # ... with 1,694 more rows
```

- How does life expectancy (lifeExp) change over time (year) for each country (country)?

```
gapminder %>%  
  ggplot(aes(year, lifeExp, group = country)) +  
  geom_line(alpha = 1/3)
```






Model for a single country

```
nz <- filter(gapminder, country == "New Zealand")
nz %>% ggplot(aes(year, lifeExp)) + geom_line() + ggtitle("Full data = ")

nz_mod <- lm(lifeExp ~ year, data = nz)
nz %>% add_predictions(nz_mod) %>%
  ggplot(aes(year, pred)) + geom_line() + ggtitle("Linear trend + ")

nz %>% add_residuals(nz_mod) %>%
  ggplot(aes(year, resid)) +
    geom_hline(yintercept = 0, colour = "white", size = 3) +
    geom_line() + ggtitle("Remaining pattern")
```



```
(by_country <- gapminder %>%  
  group_by(country, continent) %>%  
  nest())  
#> # A tibble: 142 x 3  
#> # Groups:   country, continent [710]  
#>   country      continent      data  
#>   <fct>        <fct>    <list<df[,4]>>  
#> 1 Afghanistan Asia      [12 x 4]  
#> 2 Albania      Europe    [12 x 4]  
#> 3 Algeria      Africa    [12 x 4]  
#> 4 Angola       Africa    [12 x 4]  
#> 5 Argentina    Americas  [12 x 4]  
#> 6 Australia    Oceania   [12 x 4]  
#> 7 Austria      Europe    [12 x 4]  
#> 8 Bahrain      Asia      [12 x 4]  
#> 9 Bangladesh   Asia      [12 x 4]  
#> 10 Belgium     Europe    [12 x 4]  
#> # ... with 132 more rows
```

- In a grouped data frame, each row is an observation.
- In a nested data frame, each row is a group.

■ A model-fitting function applied to every country:

```
country_model <- function(df) lm(lifeExp ~ year, data = df)
models <- map(by_country$data, country_model)
```

■ Or add an additional list-column:

```
(by_country <- by_country %>%
  mutate(model = map(data, country_model)))
#> # A tibble: 142 x 4
#> # Groups:   country, continent [710]
#>   country      continent      data model
#>   <fct>        <fct>    <list<df[,4]>> <list>
#> 1 Afghanistan Asia      [12 x 4] <lm>
#> 2 Albania      Europe    [12 x 4] <lm>
#> 3 Algeria      Africa    [12 x 4] <lm>
#> 4 Angola       Africa    [12 x 4] <lm>
#> 5 Argentina    Americas  [12 x 4] <lm>
#> 6 Australia    Oceania   [12 x 4] <lm>
#> 7 Austria      Europe    [12 x 4] <lm>
#> 8 Bahrain      Asia      [12 x 4] <lm>
#> 9 Bangladesh   Asia      [12 x 4] <lm>
#> 10 Belgium     Europe    [12 x 4] <lm>
#> # ... with 132 more rows
```

- Avoid leaving the list of models as a free-floating object.
- No need to manually keep them in sync when using e.g. `filter` or `arrange`.

```
by_country %>% filter(continent == "Europe")
#> # A tibble: 30 x 4
#> # Groups:   country, continent [710]
#>   country                continent      data model
#>   <fct>                  <fct>      <list<df[,4]>> <list>
#> 1 Albania                Europe      [12 x 4] <lm>
#> 2 Austria                Europe      [12 x 4] <lm>
#> 3 Belgium                Europe      [12 x 4] <lm>
#> 4 Bosnia and Herzegovina Europe      [12 x 4] <lm>
#> 5 Bulgaria                Europe      [12 x 4] <lm>
#> 6 Croatia                Europe      [12 x 4] <lm>
#> 7 Czech Republic         Europe      [12 x 4] <lm>
#> 8 Denmark                Europe      [12 x 4] <lm>
#> 9 Finland                Europe      [12 x 4] <lm>
#> 10 France                Europe      [12 x 4] <lm>
#> # ... with 20 more rows
```


Adding residuals

```
(by_country <- by_country %>%  
  mutate(resids = map2(data, model, add_residuals)))  
#> # A tibble: 142 x 5  
#> # Groups:   country, continent [710]  
#>   country      continent      data model  resids  
#>   <fct>      <fct>      <list<df[,4]>> <list> <list>  
#> 1 Afghanistan Asia      [12 x 4] <lm>   <tibble [12 x 5]>  
#> 2 Albania    Europe    [12 x 4] <lm>   <tibble [12 x 5]>  
#> 3 Algeria    Africa    [12 x 4] <lm>   <tibble [12 x 5]>  
#> 4 Angola     Africa    [12 x 4] <lm>   <tibble [12 x 5]>  
#> 5 Argentina  Americas  [12 x 4] <lm>   <tibble [12 x 5]>  
#> 6 Australia  Oceania   [12 x 4] <lm>   <tibble [12 x 5]>  
#> 7 Austria    Europe    [12 x 4] <lm>   <tibble [12 x 5]>  
#> 8 Bahrain    Asia      [12 x 4] <lm>   <tibble [12 x 5]>  
#> 9 Bangladesh Asia      [12 x 4] <lm>   <tibble [12 x 5]>  
#> 10 Belgium   Europe    [12 x 4] <lm>   <tibble [12 x 5]>  
#> # ... with 132 more rows
```

```
resids <- unnest(by_country, resids)
resids
#> # A tibble: 1,704 x 9
#> # Groups:   country, continent [710]
#>   country continent   data model  year lifeExp   pop gdpPercap
#>   <fct>    <fct>      <list<d> <lis> <int>   <dbl>   <int>    <dbl>
#> 1 Afghan~ Asia      [12 x 4] <lm>  1952   28.8 8.43e6   779.
#> 2 Afghan~ Asia      [12 x 4] <lm>  1957   30.3 9.24e6   821.
#> 3 Afghan~ Asia      [12 x 4] <lm>  1962   32.0 1.03e7   853.
#> 4 Afghan~ Asia      [12 x 4] <lm>  1967   34.0 1.15e7   836.
#> 5 Afghan~ Asia      [12 x 4] <lm>  1972   36.1 1.31e7   740.
#> 6 Afghan~ Asia      [12 x 4] <lm>  1977   38.4 1.49e7   786.
#> 7 Afghan~ Asia      [12 x 4] <lm>  1982   39.9 1.29e7   978.
#> 8 Afghan~ Asia      [12 x 4] <lm>  1987   40.8 1.39e7   852.
#> 9 Afghan~ Asia      [12 x 4] <lm>  1992   41.7 1.63e7   649.
#> 10 Afghan~ Asia     [12 x 4] <lm>  1997   41.8 2.22e7   635.
#> # ... with 1,694 more rows, and 1 more variable: resid <dbl>
```

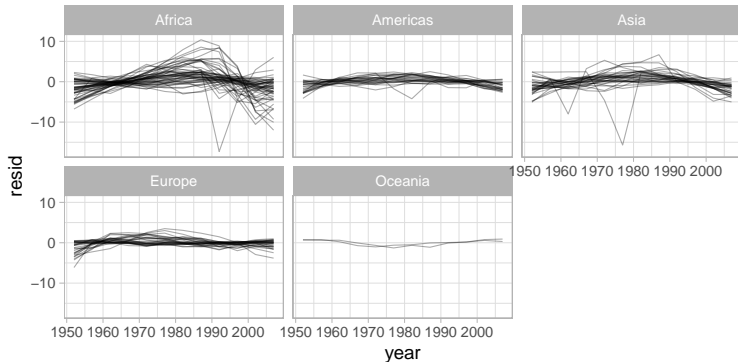
Visualizing the residuals

```
resids %>%  
  ggplot(aes(year, resid)) +  
    geom_line(aes(group = country), alpha = 1 / 3) +  
    geom_smooth(se = FALSE)
```



Visualizing the residuals cont'd

```
resids %>%  
  ggplot(aes(year, resid, group = country)) +  
    geom_line(alpha = 1 / 3) +  
    facet_wrap(~continent)
```



```
library(broom)
glance(nz_mod)
#> # A tibble: 1 x 11
#>   r.squared adj.r.squared sigma statistic p.value    df logLik   AIC
#>   <dbl>      <dbl> <dbl>      <dbl>   <dbl> <int> <dbl> <dbl>
#> 1      0.954      0.949 0.804      205. 5.41e-8     2  -13.3  32.6
#> # ... with 3 more variables: BIC <dbl>, deviance <dbl>,
#> #   df.residual <int>
by_country %>%
  mutate(glance = map(model, glance)) %>%
  unnest(glance) %>%
  print(n = 3)
#> # A tibble: 142 x 16
#> # Groups:   country, continent [710]
#>   country continent   data model resid r.squared adj.r.squared
#>   <fct>   <fct>      <list<d> <lis> <list>      <dbl>      <dbl>
#> 1 Afghan~ Asia    [12 x 4] <lm>  <tibb~    0.948      0.942
#> 2 Albania Europe  [12 x 4] <lm>  <tibb~    0.911      0.902
#> 3 Algeria Africa  [12 x 4] <lm>  <tibb~    0.985      0.984
#> # ... with 139 more rows, and 9 more variables: sigma <dbl>,
#> #   statistic <dbl>, p.value <dbl>, df <int>, logLik <dbl>,
#> #   AIC <dbl>, BIC <dbl>, deviance <dbl>, df.residual <int>
```


```
by_country_glance <- by_country %>%
  mutate(glance = map(model, glance)) %>%
  unnest(glance)
by_country_glance
#> # A tibble: 142 x 16
#> # Groups:   country, continent [710]
#>   country continent      data model residr r.squared adj.r.squared
#>   <fct>    <fct>      <list<d> <lis> <list>      <dbl>      <dbl>
#> 1 Afghan~ Asia      [12 x 4] <lm>  <tibb~      0.948      0.942
#> 2 Albania Europe    [12 x 4] <lm>  <tibb~      0.911      0.902
#> 3 Algeria Africa    [12 x 4] <lm>  <tibb~      0.985      0.984
#> 4 Angola  Africa    [12 x 4] <lm>  <tibb~      0.888      0.877
#> 5 Argent~ Americas [12 x 4] <lm>  <tibb~      0.996      0.995
#> 6 Austra~ Oceania   [12 x 4] <lm>  <tibb~      0.980      0.978
#> 7 Austria Europe    [12 x 4] <lm>  <tibb~      0.992      0.991
#> 8 Bahrain Asia      [12 x 4] <lm>  <tibb~      0.967      0.963
#> 9 Bangla~ Asia      [12 x 4] <lm>  <tibb~      0.989      0.988
#> 10 Belgium Europe   [12 x 4] <lm>  <tibb~      0.995      0.994
#> # ... with 132 more rows, and 9 more variables: sigma <dbl>,
#> #   statistic <dbl>, p.value <dbl>, df <int>, logLik <dbl>,
#> #   AIC <dbl>, BIC <dbl>, deviance <dbl>, df.residual <int>
```


Which models don't fit well?

```
by_country_glance %>%  
  arrange(r.squared)  
#> # A tibble: 142 x 16  
#> # Groups:   country, continent [710]  
#>   country continent      data model resid r.squared adj.r.squared  
#>   <fct>    <fct>      <list<d> <lis> <list>      <dbl>      <dbl>  
#> 1 Rwanda  Africa    [12 x 4] <lm> <tibb~    0.0172     -0.0811  
#> 2 Botswa~ Africa    [12 x 4] <lm> <tibb~    0.0340     -0.0626  
#> 3 Zimbab~ Africa    [12 x 4] <lm> <tibb~    0.0562     -0.0381  
#> 4 Zambia  Africa    [12 x 4] <lm> <tibb~    0.0598     -0.0342  
#> 5 Swazil~ Africa    [12 x 4] <lm> <tibb~    0.0682     -0.0250  
#> 6 Lesotho Africa    [12 x 4] <lm> <tibb~    0.0849     -0.00666  
#> 7 Cote d~ Africa    [12 x 4] <lm> <tibb~    0.283      0.212  
#> 8 South ~ Africa    [12 x 4] <lm> <tibb~    0.312      0.244  
#> 9 Uganda  Africa    [12 x 4] <lm> <tibb~    0.342      0.276  
#> 10 Congo,~ Africa    [12 x 4] <lm> <tibb~    0.348      0.283  
#> # ... with 132 more rows, and 9 more variables: sigma <dbl>,  
#> #   statistic <dbl>, p.value <dbl>, df <int>, logLik <dbl>,  
#> #   AIC <dbl>, BIC <dbl>, deviance <dbl>, df.residual <int>
```

```
bad_fit <- filter(by_country_glance, r.squared < 0.25)
```

```
gapminder %>%  
  semi_join(bad_fit, by = "country") %>%  
  ggplot(aes(year, lifeExp, colour = country)) +  
  geom_line()
```



- `broom::glance(model)`
 - ▶ A row for each model.
 - ▶ Columns give a model summary (measure of model quality, complexity, or combination of both).
- `broom::tidy(model)`
 - ▶ A row for each coefficient in the model.
 - ▶ Columns give information about the estimate or its variability.
- `broom::augment(model, data)`
 - ▶ A row for each row in data.
 - ▶ Adds extra values like residuals, and influence statistics.

```
## Annette Dobson (1990) "An Introduction to Generalized Linear Models".  
## Page 9: Plant Weight Data.  
ctl <- c(4.17, 5.58, 5.18, 6.11, 4.50, 4.61, 5.17, 4.53, 5.33, 5.14)  
trt <- c(4.81, 4.17, 4.41, 3.59, 5.87, 3.83, 6.03, 4.89, 4.32, 4.69)  
group <- gl(2, 10, 20, labels = c("Ctl", "Trt"))  
weight <- c(ctl, trt)  
lm_D9 <- lm(weight ~ group)
```

- `broom::glance(model)`
 - ▶ A row for each model.
 - ▶ Columns give a model summary.

```
tidy(lm_D9)
#> # A tibble: 2 x 5
#>   term          estimate std.error statistic  p.value
#>   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
#> 1 (Intercept)    5.03      0.220    22.9 9.55e-15
#> 2 groupTrt     -0.371     0.311    -1.19 2.49e- 1
```

- `broom::tidy(model)`
 - ▶ A row for each coefficient in the model.
 - ▶ Columns give information about the estimate or its variability.

```
tidy(lm_D9)
#> # A tibble: 2 x 5
#>   term          estimate std.error statistic  p.value
#>   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
#> 1 (Intercept)    5.03      0.220    22.9 9.55e-15
#> 2 groupTrt     -0.371     0.311    -1.19 2.49e- 1
```

■ broom::augment(model, data)

- ▶ A row for each row in data.
- ▶ Adds extra values like residuals, and influence statistics.

```
augment(lm_D9) %>%  
  print(n = 10)  
#> # A tibble: 20 x 9  
#>   weight group .fitted .se.fit .resid .hat .sigma .cooksd  
#>   <dbl> <fct>   <dbl>   <dbl> <dbl> <dbl> <dbl>   <dbl>  
#> 1  4.17 Ctl     5.03   0.220 -0.862 0.10  0.682 0.0946  
#> 2  5.58 Ctl     5.03   0.220  0.548 0.10  0.703 0.0382  
#> 3  5.18 Ctl     5.03   0.220  0.148 0.1   0.716 0.00279  
#> 4  6.11 Ctl     5.03   0.220  1.08  0.1   0.661 0.148  
#> 5  4.5  Ctl     5.03   0.220 -0.532 0.1   0.704 0.0360  
#> 6  4.61 Ctl     5.03   0.220 -0.422 0.1   0.708 0.0227  
#> 7  5.17 Ctl     5.03   0.220  0.138 0.1   0.716 0.00242  
#> 8  4.53 Ctl     5.03   0.220 -0.502 0.1   0.705 0.0321  
#> 9  5.33 Ctl     5.03   0.220  0.298 0.1   0.713 0.0113  
#> 10 5.14 Ctl     5.03   0.220  0.108 0.1   0.716 0.00148  
#> # ... with 10 more rows, and 1 more variable: .std.resid <dbl>
```

- 1 Model basics
- 2 Model building
- 3 Many models
- 4 List-columns**

```
tibble(x = list(1:3, 3:5),
       y = c("1, 2", "3, 4, 5"))
#> # A tibble: 2 x 2
#>   x           y
#>   <list>      <chr>
#> 1 <int [3]> 1, 2
#> 2 <int [3]> 3, 4, 5
```

```
tribble(~x, ~y,
        1:3, "1, 2",
        3:5, "3, 4, 5")
#> # A tibble: 2 x 2
#>   x           y
#>   <list>      <chr>
#> 1 <int [3]> 1, 2
#> 2 <int [3]> 3, 4, 5
```

- Create the list-column:
 - ▶ With `nest()` to convert a grouped data frame into a nested data frame where you have list-column of data frames.
 - ▶ With `mutate()` and vectorised functions that return a list.
 - ▶ With `summarize()` and summary functions that return multiple results.
- Create other intermediate list-columns by transforming existing list columns with `map()`, `map2()` or `pmap()`.
- Simplify the list-column back down to a data frame or atomic vector.

Create with nesting I

```
gapminder %>%  
  group_by(country, continent) %>%  
  nest()  
#> # A tibble: 142 x 3  
#> # Groups:   country, continent [710]  
#>   country      continent      data  
#>   <fct>        <fct>      <list<df[,4]>>  
#> 1 Afghanistan Asia      [12 x 4]  
#> 2 Albania     Europe   [12 x 4]  
#> 3 Algeria     Africa   [12 x 4]  
#> 4 Angola      Africa   [12 x 4]  
#> 5 Argentina   Americas [12 x 4]  
#> 6 Australia   Oceania  [12 x 4]  
#> 7 Austria     Europe   [12 x 4]  
#> 8 Bahrain     Asia     [12 x 4]  
#> 9 Bangladesh  Asia     [12 x 4]  
#> 10 Belgium    Europe   [12 x 4]  
#> # ... with 132 more rows
```

Create with nesting II

```
gapminder %>%  
  nest(data = year:gdpPercap)  
#> # A tibble: 142 x 3  
#>   country      continent      data  
#>   <fct>      <fct>      <list<df[,4]>>  
#> 1 Afghanistan Asia      [12 x 4]  
#> 2 Albania    Europe    [12 x 4]  
#> 3 Algeria    Africa    [12 x 4]  
#> 4 Angola     Africa    [12 x 4]  
#> 5 Argentina  Americas  [12 x 4]  
#> 6 Australia  Oceania   [12 x 4]  
#> 7 Austria    Europe    [12 x 4]  
#> 8 Bahrain    Asia      [12 x 4]  
#> 9 Bangladesh Asia      [12 x 4]  
#> 10 Belgium   Europe    [12 x 4]  
#> # ... with 132 more rows
```

Create from vectorized functions

```
df <- tribble(~x1, "a,b,c", "d,e,f,g")
df %>%
  mutate(x2 = stringr::str_split(x1, ","))
#> # A tibble: 2 x 2
#>   x1      x2
#>   <chr>   <list>
#> 1 a,b,c   <chr [3]>
#> 2 d,e,f,g <chr [4]>

sim <- tribble(~f,      ~params,
               "runif", list(min = -1, max = -1),
               "rnorm", list(sd = 5),
               "rpois", list(lambda = 10))
sim %>%
  mutate(sims = invoke_map(f, params, n = 10))
#> # A tibble: 3 x 3
#>   f      params      sims
#>   <chr> <list>      <list>
#> 1 runif <named list [2]> <dbl [10]>
#> 2 rnorm <named list [1]> <dbl [10]>
#> 3 rpois <named list [1]> <int [10]>
```

■ What's wrong here?

```
mtcars %>%  
  group_by(cyl) %>%  
  summarize(q = quantile(mpg))  
#> Error: Column `q` must be length 1 (a summary value), not 5
```

■ Use list-columns:

```
mtcars %>%  
  group_by(cyl) %>%  
  summarize(q = list(quantile(mpg)))  
#> # A tibble: 3 x 2  
#>   cyl q  
#>   <dbl> <list>  
#> 1     4 <dbl [5]>  
#> 2     6 <dbl [5]>  
#> 3     8 <dbl [5]>
```

```
probs <- c(0.01, 0.25, 0.5, 0.75, 0.99)
mtcars %>%
  group_by(cyl) %>%
  summarize(p = list(probs),
            q = list(quantile(mpg, probs))) %>%
  unnest(cols = c(p, q)) %>%
  print(n = 7)
#> # A tibble: 15 x 3
#>   cyl      p      q
#>   <dbl> <dbl> <dbl>
#> 1     4  0.01  21.4
#> 2     4  0.25  22.8
#> 3     4  0.5   26
#> 4     4  0.75  30.4
#> 5     4  0.99  33.8
#> 6     6  0.01  17.8
#> 7     6  0.25  18.6
#> # ... with 8 more rows
```

- If you want a single value, use `mutate()` with `map_lgl()`, `map_int()`, `map_dbl()`, and `map_chr()` to create an atomic vector.
- If you want many values, use `unnest()` to convert list-columns back to regular columns, repeating the rows as many times as necessary.

```
df <- tribble(~x, letters[1:5], 1:3, runif(5))
```

```
df %>%  
  mutate(type = map_chr(x, typeof),  
         length = map_int(x, length))
```

```
#> # A tibble: 3 x 3  
#>   x           type      length  
#>   <list>    <chr>    <int>  
#> 1 <chr [5]> character      5  
#> 2 <int [3]> integer         3  
#> 3 <dbl [5]> double          5
```

- Columns with the same number of elements:

```
tibble(x = 1:2,
       y = list(1:4, 1)) %>%
  unnest(y)
#> # A tibble: 5 x 2
#>       x     y
#>   <int> <dbl>
#> 1     1     1
#> 2     1     2
#> 3     1     3
#> 4     1     4
#> 5     2     1
```

- Columns with different number of elements:

```
tribble(~x, ~y, ~z,
        1, "a", 1:2,
        2, c("b", "c"), 3
        ) %>%
  unnest(c(y, z))
#> # A tibble: 4 x 3
#>       x y     z
#>   <dbl> <chr> <dbl>
#> 1     1 a     1
#> 2     1 a     2
#> 3     2 b     3
#> 4     2 c     3
```